

# COMP 3430

Operating Systems

June 5<sup>th</sup>, 2019

# Goals

By the end of today's lecture, you should be able to:

- Describe scheduling policies
- Describe preemptive multitasking
- Evaluate scheduling policies
- Write a UNIX program that implements scheduling policies



© S A Kindstrom CC BY-NC 2.0

# This week *and* next week

Let's take a look at the schedule.



© I am CC BY-NC-ND 2.0

# Keep, start, stop.

*Your opinion is very important to us.  
We appreciate your feedback and will  
use it to evaluate changes and make  
improvements in our site COURSE.*



The tags for this image are  
“Excited Person Happy Young Woman Joy Happiness”  
(Pixabay License)

# Scheduling

- Processes have a lifecycle.
  - They have *states*.
- We've got all these processes/threads talking to each other.
  - Scheduling: decide which one should use resources *next*.



How are we supposed to schedule *that*? (Pixabay License)

# Scheduling goals

- Schedulers have different (often competing) goals:
  - Minimizing turnaround time ( $T_{\text{completion}} - T_{\text{arrival}}$ ),
  - Maximizing throughput (get as much work done as possible),
  - Being fair to all work that needs to be done (nobody is starved of access to resources),
  - Minimizing response time ( $T_{\text{firstrun}} - T_{\text{arrival}}$ ).

# Machinery

Two main parts to the scheduling parts of an OS:

1. The **scheduler** (an oracle that decides which job goes next)
2. The **dispatcher** (the machine that actually *switches* to that job)

Our focus is on the **scheduler** and the algorithms it might use.

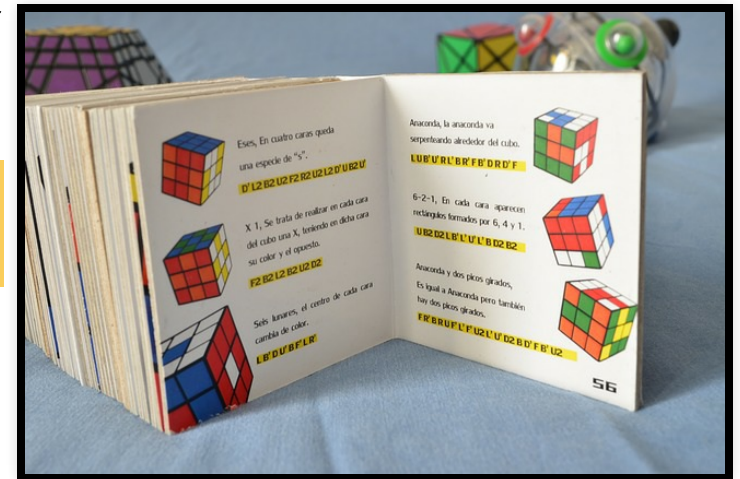


Pixabay License

# Algorithms

With the person beside you (or thinking about it by yourself, don't be pushy):

Remind yourself {f,ves}: which scheduling algorithms did we look at in the first class?



Pixabay License



# Weaknesses

Discussion: What kind of *weaknesses* do these algorithms have? **Why are they bad?**

Think about how each of these schedulers might perform with respect to the goals of scheduling.

1. First come first serve/first in first out(FIFO):

- Minimizing turnaround time: not good with turnaround time.
- Maximizing throughput: good
- Fair to all work: bad(if the first job never get finished, last job cant get access to resources.)
- Minimizing response time: not good.(same reason as last one.)

2. Shortest remaining time first:

pick the smallest amount of remaining resource utilization.


- Minimizing turnaround time: not good with turnaround time.
- Maximizing throughput: good
- Fair to all work: bad(not fair at all)
- Minimizing response time: not good.(same reason as last one.)

3. priority scheduling:

who gets to decide what priorities are?

-this will have the same problems with SRTF.

None them are really good.



© Zebeth

# Preemption

- **Main** problem: have to wait for a job to complete before scheduling next.
- Idea: Let the OS **stop** a running job *execution*.
  - When the OS regains control of the processor it can *schedule* then *dispatch* a new job.



Pixabay License

# When?

Think about it: When might the OS have an opportunity to regain control of the processor?

# Now, that's when

- The OS can regain control at any time (...with *hardware assistance*).
  - The **hardware** can (partially) suspend a job and call code in the kernel.
  - This *interruption* usually happens at specified time intervals (literally a clock ).

# Round-robin

- **Idea:** Let each job have some set period of time on the processor ( $n$  milliseconds)
  - After that time period has elapsed, *switch* to a different job.
- Let's tie some shoes .



Pixabay License

# Problems?

Discuss (or think): What problems might round-robin have? Consider each of the metrics that we defined earlier:

- Turnaround time ( $T_{\text{completion}} - T_{\text{arrival}}$ ),
- Maximizing throughput (get as much work done as possible),
- Being fair to all work that needs to be done (nobody is starved of access to resources),
- Minimizing response time ( $T_{\text{firstrun}} - T_{\text{arrival}}$ ).

Round-Robin Algorithm:

- Turnaround time: Bad
- Maximizing throughput: good
- Being fair: excellent
- Minimizing response time: good,  
(better than FIFO)

# What's the kernel doing?

Let's (briefly) take a guided tour through the Linux kernel source code.

Start with <https://elixir.bootlin.com/linux/latest/source/kernel/sched/core.c>,  
go to line 3555.

```
asmlinkage __visible void __sched schedule(void)
```

- What is `asmlinkage`?
- What is `__visible`?

# Goals

By the end of today's lecture, you should be able to:

- Describe scheduling policies (we have seen several)
- Describe preemptive multitasking (how does the OS *do* this?)
- Evaluate scheduling policies (optimizing for *metrics*)
- Write a UNIX program that implements scheduling policies ?



© S A Kindstrom CC BY-NC 2.0



